

Week 1

Reporting Fire Drill

Turn a recurring KPI scramble into a reproducible automation.

BUILD

BREAK

DEFEND

A mission-based alternative course design

Why this week matters

The practical motivation

- Most analytics teams do not first need a grand AI strategy; they need fewer manual fire drills.
- The agent is useful when it turns repeated analyst labour into a checked workflow.
- The risk is metric drift: the report looks polished while the definition quietly changed.

Monday, 8:40 a.m. The COO wants the weekly KPI pack by 10:00. Last week the analyst spent three hours copying numbers between CSV files, Excel, and slides.

The studio rhythm

1

BUILD

Use the agent to create the first useful version.

2

BREAK

Attack the draft with bad data, edge cases, and misuse.

3

DEFEND

Keep only what survives evidence, tests, and explanation.

Mission brief

Use an AI agent to build a reporting workflow that can be rerun, checked, and defended.

You will build

Repo with work order, metric definitions, reporting script, smoke tests, and AI-use log.

Success looks like

- The artefact runs from a clean checkout.
- At least one agent failure is found and fixed.
- The final claim has evidence a peer can inspect.

Start here

Use these files and assumptions before asking the agent to build.

Starter materials

- Starter_Kit/sample_data/week1_customers.csv
- Starter_Kit/sample_data/week1_sales.csv and week1_support.csv
- Starter_Kit/sample_data/week1_metric_definitions.md
- Starter_Kit/templates/work_order_template.md and ai_use_log_template.md

Confirm first

- Which definition of active customer is approved?
- Which known totals should the report match?
- Where will the rejected agent output be logged?

Worked example to imitate

Use Worked_Examples.md and Starter_Kit/examples/ for the full artefact.

What strong evidence looks like

- Good work order: names files, definitions, constraints, and known checks.
- Good AI-use log: accepts file-loading structure but rejects recent-login denominator.
- Good evidence: paid active = 6, net revenue = 2660, open tickets = 3.

Discuss before building

- What makes this example inspectable?
- Which claim is supported by evidence?
- What limit is still named?

Concept burst

Only the ideas needed for today's build

1

Agent as junior analyst with tools, not as oracle.

2

Metric definition as a contract.

3

Smoke tests before stakeholder confidence.

Guided practice menu

Use the matching file in Exercises/ for timing, instructions, and instructor notes.

Core exercises

- Metric contract clinic: define each KPI before prompting.
- Weak-prompt rewrite: turn 'make a KPI report' into a work order.
- Smoke-test hunt: create checks that fail when the denominator changes.

Submit by the end

- One concrete artefact.
- One reproduced or plausible failure.
- One evidence-linked defence note.

Three-hour delivery path

Use these slides as pacing support: set context, practise, build, break, repair, defend.

Week 1

180-minute run of show

- 1 0-5: retrieval link and setup check.
- 2 5-15: cold open and stakeholder pressure.
- 3 15-30: demo target and success criteria.
- 4 30-50: concept burst with worked example.
- 5 50-80: guided practice and prompt clinic.
- 6 80-115: Sprint A builds the first version.
- 7 115-135: failure cards and defect exposure.
- 8 135-165: Sprint B repairs and evidence pack.
- 9 165-180: peer demo, defence note, exit ticket.

Learning outcomes for today

Build

Create the first version of: Repo with work order, metric definitions, reporting script, smoke tests, and AI-use log.

Break

Expose at least one agent, data, prompt, model, or workflow failure.

Defend

Explain what can be trusted, what was rejected, and what still needs work.

Team roles

Driver

Operates the repo, files, agent, notebook, tool, or spreadsheet.

Reviewer

Reads outputs, checks assumptions, challenges claims, and watches for defects.

Reporter

Maintains the evidence pack, AI-use log, and defence note.

Retrieval link

Ask the room

- What did we build, break, and defend last week?
- Which habit matters again today?
- Which previous artefact becomes input for this mission?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Cold open discussion

Ask the room

- Monday, 8:40 a.m. The COO wants the weekly KPI pack by 10:00. Last week the analyst spent three hours copying numbers between CSV files, Excel, and slides.
- What would go wrong if the team trusted the agent too early?
- What evidence would make the stakeholder comfortable?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Stakeholder lens

1 Name the stakeholder and the decision they must make.

2 Name the artefact they expect to receive.

3 Name the risk if the artefact is wrong.

4 Name the evidence that would reduce that risk.

5 Name what should be out of scope today.

What makes this an AI-agent course?

Delegation

Students delegate bounded work to the agent instead of doing everything manually.

Supervision

Students inspect, test, reject, and repair agent output.

Accountability

Students defend the final artefact as human analysts, not as agent operators.

Understand the problem before prompting

Students should know the files, definitions, stakeholder, and acceptance criteria before the agent starts.

Week 1

Starter file walkthrough

- Starter_Kit/sample_data/week1_customers.csv
- Starter_Kit/sample_data/week1_sales.csv and week1_support.csv
- Starter_Kit/sample_data/week1_metric_definitions.md
- Starter_Kit/templates/work_order_template.md and ai_use_log_template.md

Confirm before prompting

Ask the room

- Which definition of active customer is approved?
- Which known totals should the report match?
- Where will the rejected agent output be logged?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Good artefact, weak artefact

Good

Runs or can be inspected, names assumptions, and includes checks.

Weak

Looks polished but hides definitions, dropped records, leakage, or unsupported claims.

Repairable

Imperfect first draft with visible evidence and a clear next fix.

Control the agent before it controls the task

Students practise turning vague requests into supervised work orders.

Week 1

Weak prompt pattern

Too broad

The agent chooses the task boundary and success criteria.

Too trusting

The agent chooses definitions, joins, thresholds, or causal claims.

Too vague

The output cannot be checked because no acceptance criteria were given.

Strong work-order pattern

- 1 Business goal: what decision or workflow does this support?
- 2 Inputs: files, tables, definitions, sources, or prior artefacts.
- 3 Task: the bounded work the agent should perform.
- 4 Acceptance criteria: what must be true before humans accept it.
- 5 Constraints: what the agent must not invent, drop, infer, or overclaim.
- 6 Output format: code, table, memo, test, checklist, or tool.

Activity 1 - prompt rewrite

10 min

Metric contract clinic: define each KPI before prompting.

Output before moving on

- One revised work order.
- One acceptance criterion.
- One constraint that prevents a known failure.

Prompt rewrite debrief

Debrief prompts

- Which word in the prompt reduced ambiguity?
- Which acceptance criterion can actually fail?
- Which constraint protects the stakeholder?

Instructor listens for

- A specific artefact, not a general intention.
- A named failure mode.
- Evidence that can be inspected.
- A limit or no-use condition.

Concept Depth

Short teaching, immediate use

Use these concept slides to equip the build, not to turn the seminar into a lecture.

Week 1

Concept 1: Agent as junior analyst with tools, not as oracle.

Meaning

How this idea changes today's artefact: Repo with work order, metric definitions, reporting script, smoke tests, and AI-use log.

Failure prevented

Connect it to one failure card: Metric ambiguity: active customer means any login vs paid account.

Evidence

Students must show:
The work order the agent followed.

Concept 1 check

Ask the room

- What would the agent do if this concept were ignored?
- What would a strong student artefact include?
- What is the smallest evidence item that proves the concept was used?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Concept 2: Metric definition as a contract.

Meaning

How this idea changes today's artefact: Repo with work order, metric definitions, reporting script, smoke tests, and AI-use log.

Failure prevented

Connect it to one failure card: Denominator switch: churn uses opening base in one slide and closing base in another.

Evidence

Students must show:
The metric-definition note.

Concept 2 check

Ask the room

- What would the agent do if this concept were ignored?
- What would a strong student artefact include?
- What is the smallest evidence item that proves the concept was used?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Concept 3: Smoke tests before stakeholder confidence.

Meaning

How this idea changes today's artefact: Repo with work order, metric definitions, reporting script, smoke tests, and AI-use log.

Failure prevented

Connect it to one failure card: Polished output: the chart looks good but the smoke test is missing.

Evidence

Students must show:
The red-to-green smoke-test output.

Concept 3 check

Ask the room

- What would the agent do if this concept were ignored?
- What would a strong student artefact include?
- What is the smallest evidence item that proves the concept was used?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Worked Example

Model the evidence standard

Students should imitate the structure of evidence, not copy the example.

Week 1

Worked example highlights

- Good work order: names files, definitions, constraints, and known checks.
- Good AI-use log: accepts file-loading structure but rejects recent-login denominator.
- Good evidence: paid active = 6, net revenue = 2660, open tickets = 3.

Worked example critique

Ask the room

- What claim is being made?
- Where is the evidence?
- What limit is named?
- What would you ask before accepting it?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Make students do the thinking before the build

The goal is not coverage. The goal is a better first sprint.

Week 1

Exercise 1

12 min

Metric contract clinic: define each KPI before prompting.

Output before moving on

- A visible table, prompt, test, note, or decision.
- One assumption written down.
- One peer or instructor question answered.

Exercise 1 debrief

Debrief prompts

- What was easy for the agent?
- What decision had to stay human?
- What evidence changed the team's mind?

Instructor listens for

- A specific artefact, not a general intention.
- A named failure mode.
- Evidence that can be inspected.
- A limit or no-use condition.

Exercise 2

12 min

Weak-prompt rewrite: turn 'make a KPI report' into a work order.

Output before moving on

- A visible table, prompt, test, note, or decision.
- One assumption written down.
- One peer or instructor question answered.

Exercise 2 debrief

Debrief prompts

- What was easy for the agent?
- What decision had to stay human?
- What evidence changed the team's mind?

Instructor listens for

- A specific artefact, not a general intention.
- A named failure mode.
- Evidence that can be inspected.
- A limit or no-use condition.

Exercise 3

18 min

Smoke-test hunt: create checks that fail when the denominator changes.

Output before moving on

- A visible table, prompt, test, note, or decision.
- One assumption written down.
- One peer or instructor question answered.

Exercise 3 debrief

Debrief prompts

- What was easy for the agent?
- What decision had to stay human?
- What evidence changed the team's mind?

Instructor listens for

- A specific artefact, not a general intention.
- A named failure mode.
- Evidence that can be inspected.
- A limit or no-use condition.

Sprint A

Build the first version

Time-box the build so students have something real to attack.

Week 1

Sprint A operating rules

- 1 Start from the work order, not a blank chat.
- 2 Keep the agent output visible and reviewable.
- 3 Run the artefact before improving it.
- 4 Write down the first defect or uncertainty.
- 5 Commit or save the first version before repair.

Sprint A task list

- Create repo structure and metric_definitions.md.
- Ask the agent to scaffold the reporting script.
- Run the script and capture the first failed check.

Sprint A checkpoint

Ask the room

- What runs or can be inspected right now?
- What did the agent assume?
- What evidence is missing?
- What is the first defect to investigate?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Instructor roaming script

- 1 Ask: show me the artefact.
- 2 Ask: show me what the agent changed.
- 3 Ask: show me one check or known value.
- 4 Ask: what would make this unsafe?
- 5 Do not solve unless the team is blocked by setup rather than thinking.

Break Moment

Failure is the teaching device

The class now attacks the first draft so students learn what not to trust.

Week 1

Failure card 1

Defect

Metric ambiguity: active customer means any login vs paid account.

How to expose it

Create a check, adversarial example, missing-value case, source comparison, or peer question.

What to record

Save the failing output before repairing it. This becomes portfolio evidence.

Failure card 2

Defect

Denominator switch:
churn uses opening
base in one slide and
closing base in another.

How to expose it

Create a check,
adversarial example,
missing-value case,
source comparison, or
peer question.

What to record

Save the failing output
before repairing it. This
becomes portfolio
evidence.

Failure card 3

Defect

Polished output: the chart looks good but the smoke test is missing.

How to expose it

Create a check, adversarial example, missing-value case, source comparison, or peer question.

What to record

Save the failing output before repairing it. This becomes portfolio evidence.

Defect wall

Ask the room

- Post one failure your team found.
- Mark whether it was agent, data, prompt, model, tool, or claim failure.
- Write the evidence that exposed it.

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Sprint B

Repair, harden, and produce evidence

Students now move from one-off fixing to recurrence prevention.

Week 1

Repair moves

- Pin metric definitions in markdown.
- Write three smoke tests for known totals.
- Commit the rejected agent output and the corrected version.

Evidence design checklist

- 1 Name the claim the artefact makes.
- 2 Name the failure that could make the claim false.
- 3 Show the check, test, trace, source, or review that exposes the failure.
- 4 Show what changed after repair.
- 5 State the remaining limit.

Sprint B checkpoint

Ask the room

- Can you reproduce the failure?
- Can you show the fix?
- Would the failure be caught earlier next time?
- What will you put in the evidence pack?

Capture on board

- Definitions students rely on.
- Assumptions that need evidence.
- Risks that should become tests.
- Questions the agent should not decide alone.

Make evidence social

A peer should be able to inspect the artefact and challenge the defence.

Week 1

Peer demo protocol

- 1 One minute: run or show the artefact.
- 2 One minute: show the failure found.
- 3 One minute: show the repair evidence.
- 4 One minute: peer asks what still should not be trusted.
- 5 One minute: team writes one improvement before submission.

Peer review questions

Artefact

What did you build, and can someone else inspect it?

Failure

What did the agent, data, model, tool, or answer get wrong?

Defence

Which claim is supported, and what remains out of scope?

Turn activity into assessable evidence

The defence note is where students show judgement, not just effort.

Week 1

Defence note writing

10 min

I trust this KPI pack because...

Output before moving on

- One claim the team can defend.
- One evidence item linked to that claim.
- One remaining risk or no-use condition.

Evidence pack requirements

- The work order the agent followed.
- The metric-definition note.
- The red-to-green smoke-test output.
- One AI-use-log row explaining what was accepted and rejected.

Rubric lens

Useful artefact

It runs, can be inspected, or clearly solves the mission within scope.

Agent supervision

The team shows what was accepted, rejected, changed, and verified.

Defensible judgement

Claims are specific, evidence-linked, and include limits.

Close

Make the habit portable

Students should leave with one artefact, one failure, and one defensible claim.

Week 1

Common traps to avoid

- Letting the agent define the metric.
- Accepting a beautiful chart with no known-total check.
- Back-filling the AI-use log after the fact.

Stretch options for fast teams

- Generate a one-page executive summary from the checked numbers.
- Add a scheduled run and failure notification.

After-class handoff

- 1 Commit or save the artefact.
- 2 Save the failed or rejected output.
- 3 Complete the AI-use log.
- 4 Submit the defence note.
- 5 Bring one unresolved question to the next retrieval link.

Agent work order

Give the agent a job, not a wish.

Use this structure

- Goal: reproduce the weekly KPI pack from raw exports.
- Inputs: sales.csv, customers.csv, support.csv, metric_definitions.md.
- Acceptance: script runs from a clean checkout and matches three known totals.
- Do not invent metric definitions. Ask if a denominator is unclear.

Studio sprint A - build the first version

Tasks

- Create repo structure and metric_definitions.md.
- Ask the agent to scaffold the reporting script.
- Run the script and capture the first failed check.

Instructor watchlist

- Is the prompt specific enough?
- Did the team run the artefact?
- Did the team record what the agent changed?

Failure cards

Release these after Sprint A. Students must expose the defect before fixing it.

Card 1

**Metric ambiguity:
active customer
means any login vs
paid account.**

Card 2

**Denominator
switch: churn uses
opening base in one
slide and closing
base in another.**

Card 3

**Polished output: the
chart looks good
but the smoke test
is missing.**

Studio sprint B - repair and harden

Repair moves

- Pin metric definitions in markdown.
- Write three smoke tests for known totals.
- Commit the rejected agent output and the corrected version.

Definition of done

- The failure is reproducible.
- The fix is committed.
- A future failure would be caught earlier.

Evidence pack

What students should submit

- The work order the agent followed.
- The metric-definition note.
- The red-to-green smoke-test output.
- One AI-use-log row explaining what was accepted and rejected.

Peer demo

Demo in three minutes

- Run the artefact.
- Show the failure you found.
- Show the evidence that supports the final claim.

Peer questions

- What did the agent get wrong?
- What would fail next?
- What should not be used yet?

Common traps

Watch for these

- Letting the agent define the metric.
- Accepting a beautiful chart with no known-total check.
- Back-filling the AI-use log after the fact.

If your team finishes early

Stretch options

- Generate a one-page executive summary from the checked numbers.
- Add a scheduled run and failure notification.

Exit ticket

I trust this KPI pack because...

Write one sentence. It goes into your portfolio evidence.

Leave with this

Automation is not the chart. Automation is the repeatable path from raw data to a checked decision.

BUILD / BREAK / DEFEND